

Simple document

```
library(tidyverse)

## -- Attaching packages ----- tidyverse 1.3.1 --
## v ggplot2 3.3.5      v purrr  0.3.4
## v tibble  3.1.4      v dplyr  1.0.7
## v tidyr   1.1.3      v stringr 1.4.0
## v readr   2.0.1      v forcats 0.5.1

## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()

library(caret)

## Loading required package: lattice
##
## Attaching package: 'caret'

## The following object is masked from 'package:purrr':
##
##   lift

library(visdat)
library(corrplot)

## corrplot 0.92 loaded

library(AppliedPredictiveModeling)
library(pROC)

## Type 'citation("pROC")' for a citation.
##
## Attaching package: 'pROC'

## The following objects are masked from 'package:stats':
##
##   cov, smooth, var

library(rpart.plot)

## Loading required package: rpart

library(vip)

##
## Attaching package: 'vip'

## The following object is masked from 'package:utils':
##
##   vi
```

```

library(ranger)
library(tidytext)
library(pdp)

##
## Attaching package: 'pdp'
## The following object is masked from 'package:purrr':
##
##   partial
library(lime)

##
## Attaching package: 'lime'
## The following object is masked from 'package:dplyr':
##
##   explain
ctrl <- trainControl(method = "cv",
                      summaryFunction = twoClassSummary,
                      classProbs = TRUE)

knitr::opts_chunk$set(
  fig.width = 6,
  fig.asp = .6,
  out.width = "90%"
)

```

Data pre-process

```

# Import data
dat_raw <- read.csv("./airline.csv")

# Check missing value
sapply(dat_raw, function(x) sum(is.na(x)))

```

```


##              X              Gender
##              0              0
##      customer_type              age
##              0              0
##      type_of_travel      customer_class
##              0              0
##      flight_distance      inflight_wifi_service
##              0              0
##      departure_arrival_time_convenient      ease_of_online_booking
##              0              0
##              gate_location      food_and_drink
##              0              0
##      online_boarding      seat_comfort
##              0              0
##      inflight_entertainment      onboard_service
##              0              0
##      leg_room_service      baggage_handling
##              0              0

```


```
##                checkin_service                inflight_service
##                0                0
##                cleanliness                departure_delay_in_minutes
##                0                0
##                arrival_delay_in_minutes                satisfaction
##                393                0
```

```
# data clean
dat <- dat_raw %>%
  janitor::clean_names() %>%
  select(-1) %>%
  mutate(satisfaction = recode(satisfaction,
                              "satisfied" = "yes",
                              "neutral or dissatisfied" = "no"))
```

```
# deal with missing values
deal_mis <- dat[, 21:22]
bagImp = preProcess(deal_mis, method = "bagImpute")
dat = predict(bagImp, dat)
vis_miss(deal_mis)
```



```
# sample data
set.seed(1234)
dat <- dat[sample(1:nrow(dat), 2000, replace = FALSE), ]
vis_miss(dat) ## check
```



```
# --- Split data ---
set.seed(1234)
trRow <- createDataPartition(dat$satisfaction, p = 0.8, list = F)


# Train data
train <- dat[trRow, ]
x_train <- model.matrix(satisfaction ~., train)[,-1]
y_train <- train$satisfaction

# Test data
test <- dat[-trRow, ]
x_test <- model.matrix(satisfaction ~., test)[,-1]
y_test <- test$satisfaction
```

EDA


```
# Correlation plot
corrplot(cor(x_train),
          method = "circle",
          type = "upper",
          tl.col = "black",
          tl.cex = 1.2,
          tl.srt = 70)
```

```
## Warning in corrplot(cor(x_train), method = "circle", type = "upper", tl.col =
## "black", : Not been able to calculate text margin, please try again with a clean
## new empty window using {plot.new(); dev.off()} or reduce tl.cex
```




```
# Barplot matrix for categorical variables
train %>%
  select(1:2, 4:5, 23) %>%
  pivot_longer(-5,
    names_to = "variable",
    values_to = "value") %>%
  group_by(variable, value, satisfaction) %>%
  summarize(num = n()) %>%
  ungroup() %>%
  group_by(variable, satisfaction) %>%
  mutate(percent = num / sum(num),
    indicator = case_when(value == "Eco" ~ 3,
      value == "Eco Plus" ~ 2,
      value == "Business" ~ 1,
      TRUE ~ 0)) %>%
  ggplot(aes(x = reorder_within(value, indicator, variable),
    y = percent, fill = satisfaction)) +
  geom_col(position = "dodge") +
  xlab("Barplot matrix for categorical variables") +
  coord_flip() +
  scale_x_reordered() +
  facet_wrap(~ variable, scales = "free") + theme_bw()
```

`summarise()` has grouped output by 'variable', 'value'. You can override using the `.groups` argument



```
# Density plot matrix
theme1 <- transparentTheme(trans = .5)
trellis.par.set(theme1)

plt_feature <-
  featurePlot(x = x_train[, c(-1, -2, -4, -5, -6)],
             y = as.factor(y_train),
             plot = "density",
             scales = list(x = list(relation = "free"),
                           y = list(relation = "free")),
             pch = "|", auto.key = list(columns = 2))
update(plt_feature, main = "Density Plot Matrix")
```



Model fitting

Logistic regression

```
set.seed(1234)
model.glm <- train(x = x_train,
                  y = y_train,
                  method = "glm",
                  metric = "ROC",
                  trControl = ctrl)
```

```
# Test AUC and Misclassification error rate
pred_glm_auc <- predict(model.glm, newdata = x_test, type = "prob")[,2]
roc(y_test, pred_glm_auc)$auc[1]
```

```
## Setting levels: control = no, case = yes
```

```
## Setting direction: controls < cases
```

```
## [1] 0.8923937
```

```
# Confusion matrix
```

```
pred_glm <- predict(model.glm, newdata = x_test)
confusionMatrix(data = as.factor(pred_glm),
                 reference = as.factor(y_test))
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           Reference
```

```
## Prediction  no yes
```

```
##           no  201  36
```

```
##           yes   24 139
```

```
##
```

```
##           Accuracy : 0.85
```

```
##           95% CI : (0.8112, 0.8835)
```

```
##           No Information Rate : 0.5625
```

```
##           P-Value [Acc > NIR] : <2e-16
```

```
##
```

```
##           Kappa : 0.6929
```

```
##
```

```
##           McNemar's Test P-Value : 0.1556
```

```
##
```

```
##           Sensitivity : 0.8933
```

```
##           Specificity : 0.7943
```

```
##           Pos Pred Value : 0.8481
```

```
##           Neg Pred Value : 0.8528
```

```
##           Prevalence : 0.5625
```

```
##           Detection Rate : 0.5025
```

```
##           Detection Prevalence : 0.5925
```

```
##           Balanced Accuracy : 0.8438
```

```
##
```

```
##           'Positive' Class : no
```

```
##
```


MARS

```
set.seed(1234)
model.mars <- train(x = x_train,
                   y = y_train,
                   method = "earth",
                   tuneGrid = expand.grid(degree = 1:4,
                                         nprune = 5:22),
                   metric = "ROC",
                   trControl = ctrl)

model.mars$bestTune
```

```
## nprune degree
## 47 15 3
```

```
ggplot(model.mars, highlight = T) +
  theme_bw()
```



```
## test auc and misclassification error rate
pred_mars_auc <- predict(model.mars, newdata = x_test, type = "prob")[,2]
roc(y_test, pred_mars_auc)$auc[1]
```

```
## [1] 0.9773587
```

```
pred_mars <- predict(model.mars, newdata = x_test)
pred.miserror_mars <- 1 - mean(pred_mars == y_test)
pred.miserror_mars
```

```
## [1] 0.0775
```

```
confusionMatrix(data = as.factor(pred_mars),
  reference = as.factor(y_test))
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           Reference
```

```
## Prediction  no  yes
```

```
##           no  214  20
```

```
##           yes  11 155
```

```
##
```

```
##           Accuracy : 0.9225
```

```
##           95% CI : (0.8918, 0.9467)
```

```
##           No Information Rate : 0.5625
```

```
##           P-Value [Acc > NIR] : <2e-16
```

```
##
```

```
##           Kappa : 0.8416
```

```
##
```



```
## McNemar's Test P-Value : 0.1508
##
##          Sensitivity : 0.9511
##          Specificity : 0.8857
##          Pos Pred Value : 0.9145
##          Neg Pred Value : 0.9337
##          Prevalence : 0.5625
##          Detection Rate : 0.5350
##          Detection Prevalence : 0.5850
##          Balanced Accuracy : 0.9184
##
##          'Positive' Class : no
##
```

LDA

```
set.seed(1234)
model.lda <- train(x = x_train,
                  y = y_train,
                  method = "lda",
                  metric = "ROC",
                  trControl = ctrl)

## test auc and misclassification error rate
pred_lda_auc <- predict(model.lda, newdata = x_test, type = "prob")[,2]
roc(y_test, pred_lda_auc)$auc[1]
```

```
## Setting levels: control = no, case = yes
```

```
## Setting direction: controls < cases
```

```
## [1] 0.8954413
```

```
pred_lda <- predict(model.lda, newdata = x_test)
pred.miserror_lda <- 1 - mean(pred_lda == y_test)
pred.miserror_lda
```

```
## [1] 0.1525
```

```
confusionMatrix(data = as.factor(pred_lda),
                reference = as.factor(y_test))
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##          Reference
```

```
## Prediction  no yes
```

```
##          no  201  37
```

```
##          yes  24 138
```

```
##
```

```
##          Accuracy : 0.8475
```

```
##          95% CI : (0.8085, 0.8813)
```

```
##          No Information Rate : 0.5625
```

```
##          P-Value [Acc > NIR] : <2e-16
```

```
##
```

```
##          Kappa : 0.6876
```

```
##
```

```
## McNemar's Test P-Value : 0.1244
##
##      Sensitivity : 0.8933
##      Specificity : 0.7886
##      Pos Pred Value : 0.8445
##      Neg Pred Value : 0.8519
##      Prevalence : 0.5625
##      Detection Rate : 0.5025
##      Detection Prevalence : 0.5950
##      Balanced Accuracy : 0.8410
##
##      'Positive' Class : no
##
```


Classification Tree

```
set.seed(1234)
model.tree <- train(x_train,
  y_train,
  method = "rpart",
  tuneGrid = data.frame(cp = exp(seq(-10, -6, length = 30))),
  trControl = ctrl,
  metric = "ROC")

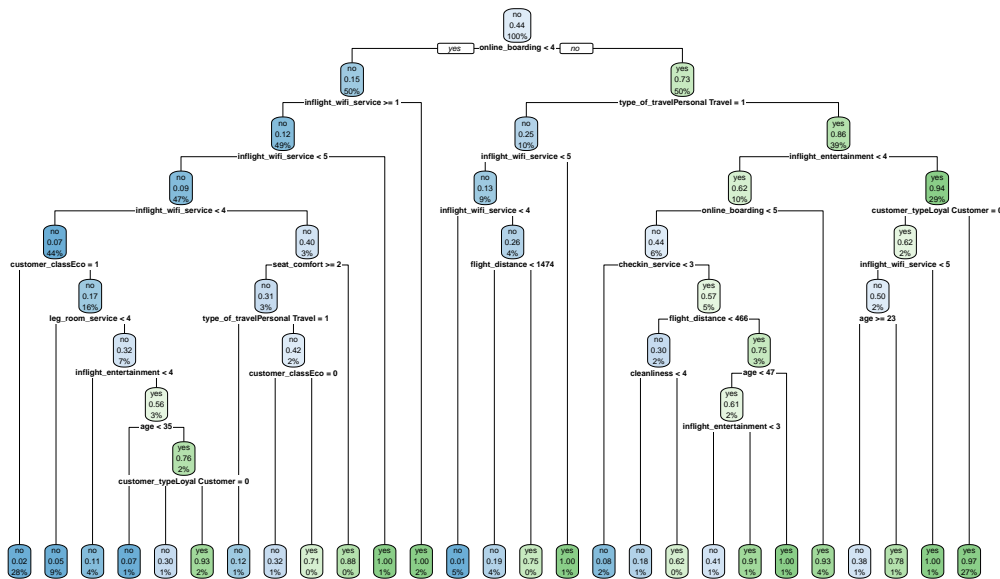
model.tree$bestTune
```

```
##      cp
## 18 0.0004735882
```

```
ggplot(model.tree, highlight = TRUE) +
  scale_x_continuous(trans = scales::log_trans(),
    breaks = scales::log_breaks()) +
  theme_bw()
```



```
rpart.plot(model.tree$finalModel)
```



```
## test auc and misclassification error rate
```

```
pred_tree_auc <- predict(model.tree, newdata = x_test, type = "prob")[,2]
roc(y_test, pred_tree_auc)$auc[1]
```

```
## [1] 0.9453333
```

```
pred_tree <- predict(model.tree, newdata = x_test)
pred.miserror_tree <- 1 - mean(pred_tree == y_test)
pred.miserror_tree
```

```
## [1] 0.1025
```

```
confusionMatrix(data = as.factor(pred_tree),
                 reference = as.factor(y_test))
```

```
## Confusion Matrix and Statistics
```

```
##
##           Reference
## Prediction no yes
##      no  201  17
##      yes  24 158
##
##           Accuracy : 0.8975
##           95% CI : (0.8635, 0.9254)
##      No Information Rate : 0.5625
##      P-Value [Acc > NIR] : <2e-16
##
##           Kappa : 0.7927
##
##  Mcnemar's Test P-Value : 0.3487
##
##           Sensitivity : 0.8933
##           Specificity : 0.9029
##      Pos Pred Value : 0.9220
##      Neg Pred Value : 0.8681
```


```
##           Prevalence : 0.5625
##           Detection Rate : 0.5025
##           Detection Prevalence : 0.5450
##           Balanced Accuracy : 0.8981
##
##           'Positive' Class : no
##
```

Random forests

```
set.seed(1234)
model.rf = train(x_train,
                 y_train,
                 method = "ranger",
                 tuneGrid = expand.grid(mtry = 3:14,
                                       splitrule = "gini",
                                       min.node.size = seq(4, 14, by = 2)),
                 metric = "ROC",
                 trControl = ctrl)

model.rf$bestTune

##      mtry splitrule min.node.size
## 49      11      gini              4
ggplot(model.rf, highlight = T) +
  theme_bw()
```



```
## test auc and misclassification error rate
pred_rf_auc <- predict(model.rf, newdata = x_test, type = "prob")[,2]
roc(y_test, pred_rf_auc)$auc[1]
```

```
## Setting levels: control = no, case = yes
```

```

## Setting direction: controls < cases
## [1] 0.9715175
pred_rf <- predict(model.rf, newdata = x_test)
pred.miserror_rf <- 1 - mean(pred_rf == y_test)
pred.miserror_rf

## [1] 0.075
confusionMatrix(data = as.factor(pred_rf),
                 reference = as.factor(y_test))

## Confusion Matrix and Statistics
##
##           Reference
## Prediction  no yes
##      no    213  18
##      yes    12 157
##
##              Accuracy : 0.925
##              95% CI : (0.8947, 0.9488)
##      No Information Rate : 0.5625
##      P-Value [Acc > NIR] : <2e-16
##
##              Kappa : 0.847
##
##  Mcnemar's Test P-Value : 0.3613
##
##      Sensitivity : 0.9467
##      Specificity : 0.8971
##      Pos Pred Value : 0.9221
##      Neg Pred Value : 0.9290
##      Prevalence : 0.5625
##      Detection Rate : 0.5325
##      Detection Prevalence : 0.5775
##      Balanced Accuracy : 0.9219
##
##      'Positive' Class : no
##

```

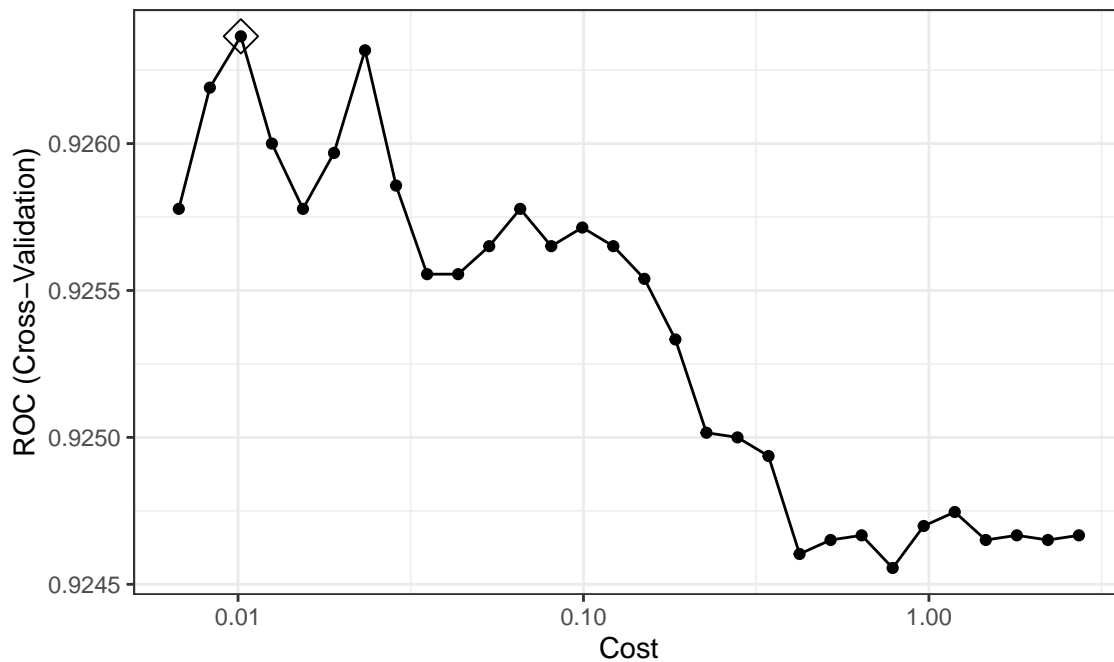
Fit a support vector classifier (linear kernel)

```

set.seed(1234)
model.svm1 = train(x_train,
                  y_train,
                  method = "svmLinear",
                  metric = "ROC",
                  tuneGrid = data.frame(C = exp(seq(-5, 1, length = 30))),
                  trControl = ctrl)

ggplot(model.svm1, highlight = TRUE) +
  scale_x_continuous(trans = scales::log_trans(),
                    breaks = scales::log_breaks()) +
  theme_bw()

```



```
## test auc and misclassification error rate
pred_svml_auc <- predict(model.svml, newdata = x_test, type = "prob")[,2]
roc(y_test, pred_svml_auc)$auc[1]
```

```
## [1] 0.8928508
```

```
pred_svml <- predict(model.svml, newdata = x_test)
pred.miserror_svml <- 1 - mean(pred_svml == y_test)
pred.miserror_svml
```

```
## [1] 0.1475
```

```
confusionMatrix(data = as.factor(pred_svml),
                 reference = as.factor(y_test))
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           Reference
```

```
## Prediction no yes
```

```
##           no  201  35
```

```
##           yes  24 140
```

```
##
```

```
##           Accuracy : 0.8525
```

```
##           95% CI : (0.8139, 0.8858)
```

```
##           No Information Rate : 0.5625
```

```
##           P-Value [Acc > NIR] : <2e-16
```

```
##
```

```
##           Kappa : 0.6982
```

```
##
```

```
##           McNemar's Test P-Value : 0.193
```

```
##
```

```
##           Sensitivity : 0.8933
```

```
##           Specificity : 0.8000
```

```
##           Pos Pred Value : 0.8517
```

```
##           Neg Pred Value : 0.8537
```

```
##           Prevalence : 0.5625
##           Detection Rate : 0.5025
##           Detection Prevalence : 0.5900
##           Balanced Accuracy : 0.8467
##
##           'Positive' Class : no
##
```

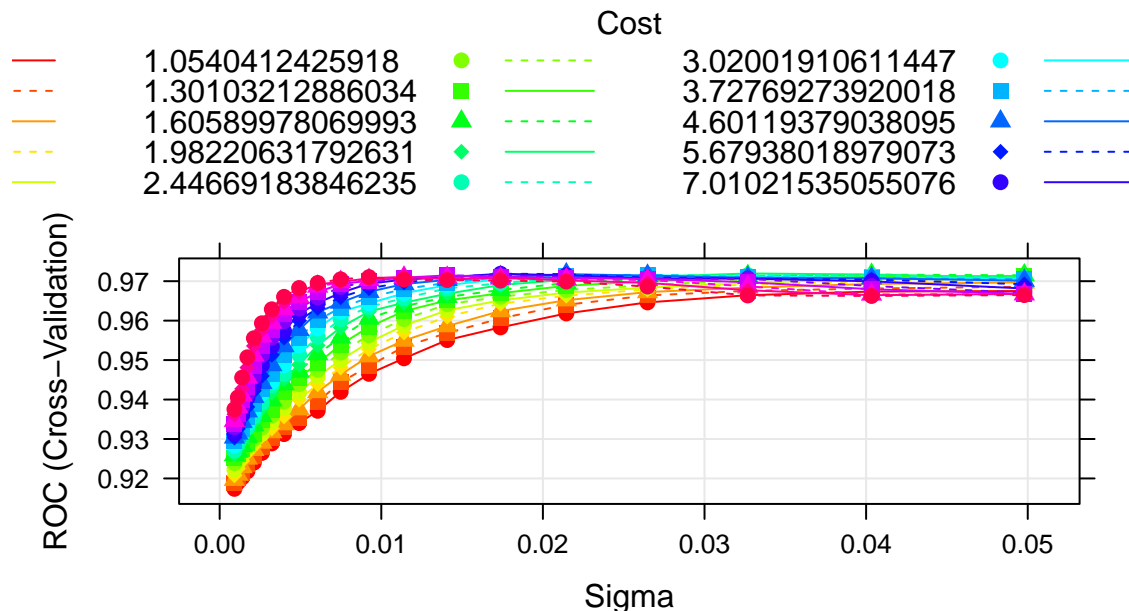
Fit a support vector machine with a radial kernel

```
set.seed(1234)
model.svmr = train(x_train,
  y_train,
  method = "svmRadialSigma",
  metric = "ROC",
  tuneGrid = expand.grid(C = exp(seq(-1, 3, length = 20)),
    sigma = exp(seq(-7, -3, length = 20))),
  trControl = ctrl)

model.svmr$bestTune
```

```
##           sigma           C
## 178 0.03267802 1.982206
```

```
myCol<- rainbow(20)
myPar <- list(superpose.symbol = list(col = myCol),
  superpose.line = list(col = myCol))
plot(model.svmr, highlight = TRUE, par.settings = myPar)
```



```
## test auc and misclassification error rate
pred_svmr_auc <- predict(model.svmr, newdata = x_test, type = "prob")[,2]
roc(y_test, pred_svmr_auc)$auc[1]
```

```
## Setting levels: control = no, case = yes
```

```
## Setting direction: controls < cases
```

```
## [1] 0.9647492
pred_svmr <- predict(model.svmr, newdata = x_test)
pred.miserror_svmr <- 1 - mean(pred_svmr == y_test)
pred.miserror_svmr
```

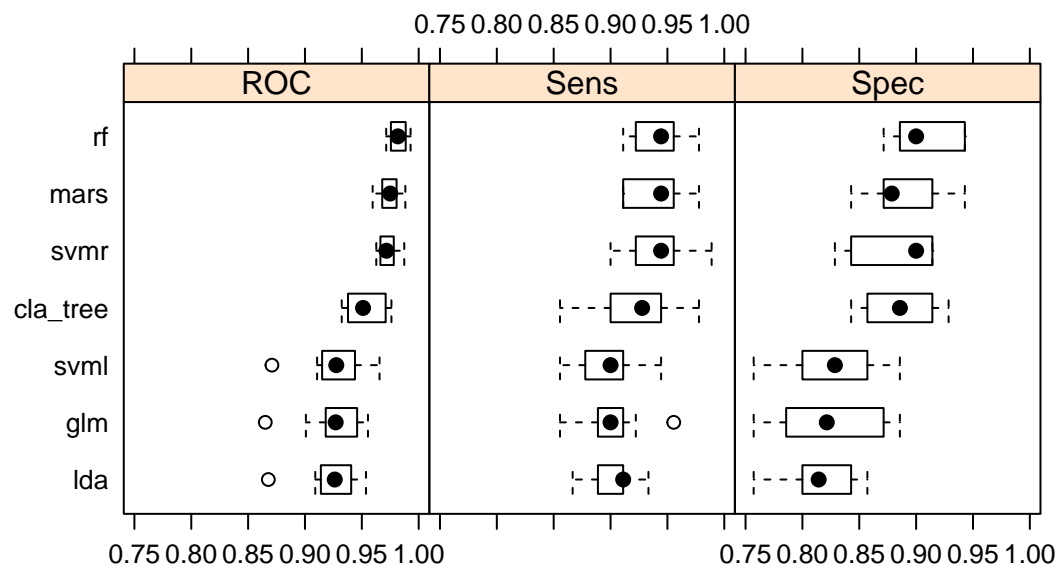
```
## [1] 0.1
confusionMatrix(data = as.factor(pred_svmr),
                 reference = as.factor(y_test))
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  no  yes
##           no 208  23
##           yes  17 152
##
##           Accuracy : 0.9
##           95% CI : (0.8663, 0.9276)
##           No Information Rate : 0.5625
##           P-Value [Acc > NIR] : <2e-16
##
##           Kappa : 0.796
##
##           Mcnemar's Test P-Value : 0.4292
##
##           Sensitivity : 0.9244
##           Specificity : 0.8686
##           Pos Pred Value : 0.9004
##           Neg Pred Value : 0.8994
##           Prevalence : 0.5625
##           Detection Rate : 0.5200
##           Detection Prevalence : 0.5775
##           Balanced Accuracy : 0.8965
##
##           'Positive' Class : no
##
```

Resample

```
resamp <- resamples(list(glm = model.glm, mars = model.mars,
                        lda = model.lda, cla_tree = model.tree,
                        rf = model.rf, svm1 = model.svm1,
                        svmr = model.svmr))


bwplot(resamp)
```

Select the rf model and interpret

```
## importance variable
set.seed(1234)
rf2.final.per <- ranger(factor(satisfaction) ~ .,
  data = train,
  mtry = model.rf$bestTune[[1]],
  min.node.size = model.rf$bestTune[[3]],
  splitrule = "gini",
  importance = "permutation",
  scale.permutation.importance = TRUE)

barplot(sort(ranger::importance(rf2.final.per), decreasing = FALSE),
  las = 2,
  horiz = TRUE,
  cex.names = 0.4,
  col = colorRampPalette(colors = c("cyan", "blue"))(8))
```



```
pdp1.rf <- model.rf %>%
  partial(pred.var = c("inflight_wifi_service")) %>%
  autoplot(train = train, rug = TRUE)

pdp2.rf <- model.rf %>%
  partial(pred.var = c("inflight_wifi_service", "age"), chull = TRUE) %>%
  autoplot(train = train, rug = TRUE)

grid.arrange(pdp1.rf, pdp2.rf, nrow = 1)
```

```
## Warning: Use of `object[[1L]]` is discouraged. Use `.`data[[1L]]` instead.
```

```
## Warning: Use of `object[["yhat"]}` is discouraged. Use `.data[["yhat"]}`
## instead.

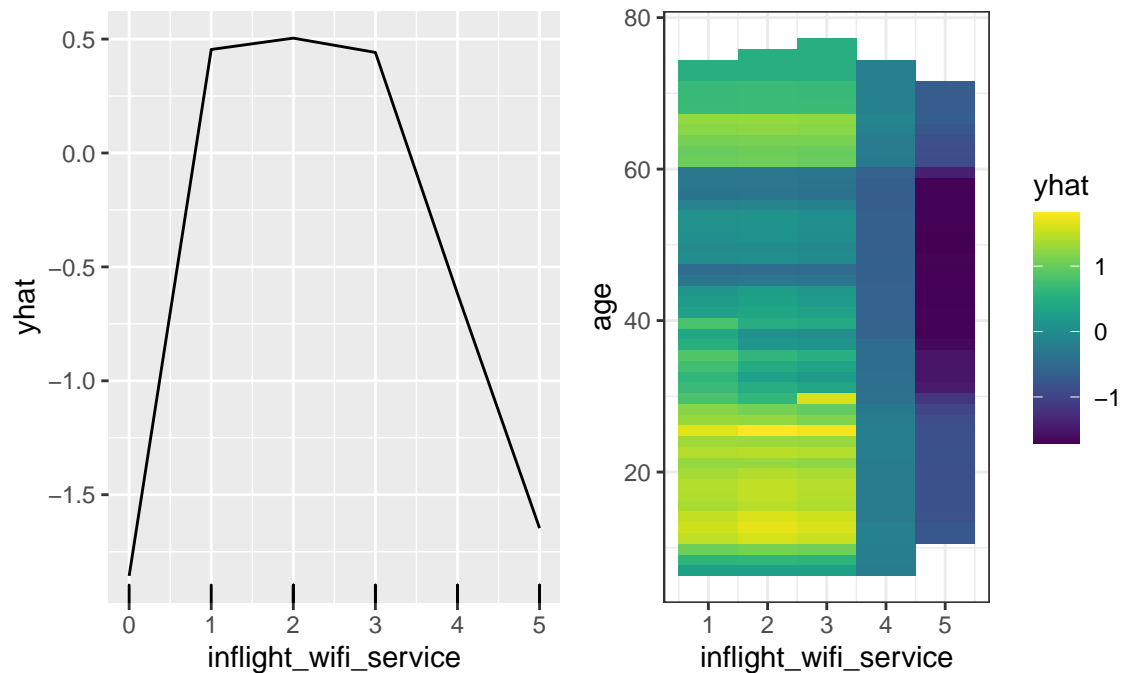
## Warning: Use of `x.rug[[1L]]` is discouraged. Use `.data[[1L]]` instead.

## Warning: Use of `object[[1L]]` is discouraged. Use `.data[[1L]]` instead.

## Warning: Use of `object[[2L]]` is discouraged. Use `.data[[2L]]` instead.

## Warning: Use of `object[["yhat"]}` is discouraged. Use `.data[["yhat"]}`
## instead.

## Warning: Use of `object[["yhat"]}` is discouraged. Use `.data[["yhat"]}`
## instead.
```



```
ice.rf <- model.rf %>%
  partial(pred.var = "inflight_wifi_service",
    grid.resolution = 100,
    ice = TRUE) %>%
  autoplot(train = dat, alpha = .1,
    center = TRUE)
```

```
## Warning: `fun.y` is deprecated. Use `fun` instead.
```

```
ice.rf
```

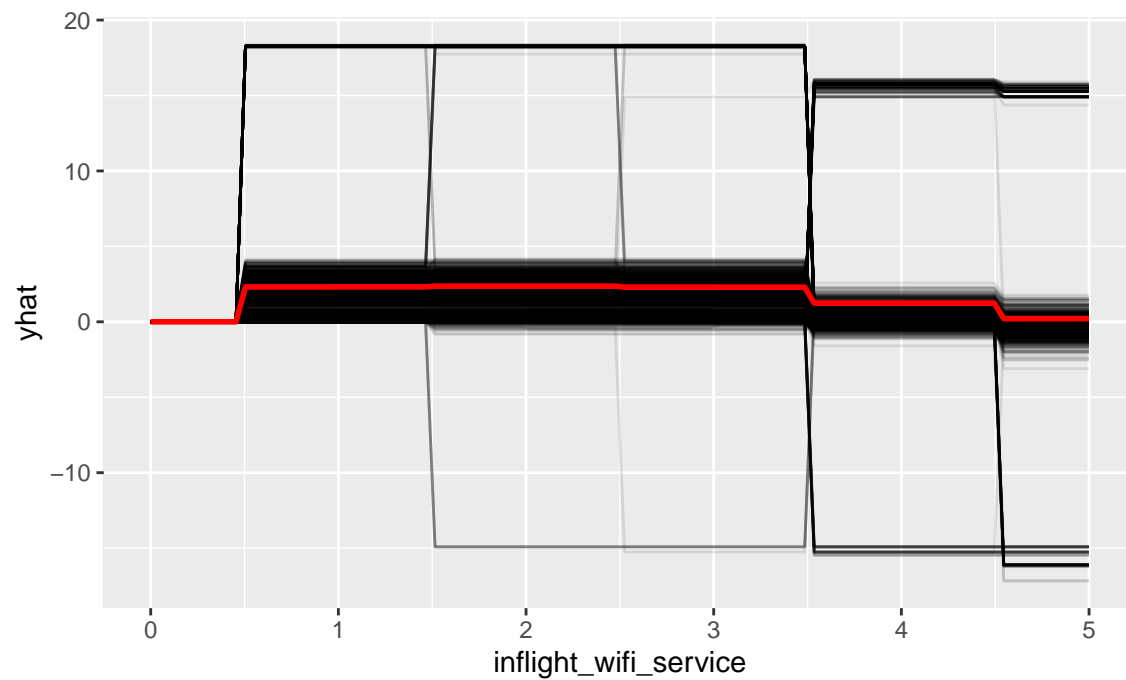
```
## Warning: Use of `object[["yhat.id"]}` is discouraged. Use `.data[["yhat.id"]}`
## instead.

## Warning: Use of `object[[1L]]` is discouraged. Use `.data[[1L]]` instead.

## Warning: Use of `object[["yhat"]}` is discouraged. Use `.data[["yhat"]}`
## instead.

## Warning: Use of `object[[1L]]` is discouraged. Use `.data[[1L]]` instead.

## Warning: Use of `object[["yhat"]}` is discouraged. Use `.data[["yhat"]}`
## instead.
```



```
##lime
#explain.rf <- lime(data.frame(x_train), model.rf)

#new_obs = x_train[1:10,]
#explana.obs = explain(data.frame(new_obs),
#                        explain.rf,
#                        n_features = 10)

#plot_features(explana.obs)
```